

CSCU9YQ - NoSQL Databases

Lecture 3: Distribution Models

Gabriela Ochoa

Outline

- Recap
 - Characteristics of NoSQL DBs
 - Data models
 - The need to scale
- Distribution models
 - Replication
 - Sharding
 - Combining Sharding and replication

Characteristics of NoSQL Databases

- Not using the Relational Model
- Running well on clusters
- Open-source
- Flexible schemas (freely add fields to DB records, without needing to define a fixed schema first)
- Big data, web applications

Data Models

- An abstract representation organising elements of data and how they relate to one another and to properties of real world entities
- The dominant data model for several decades has been the Relational Model
- The primary way in which non-tabular databases differ from relational databases is the data model
- Three main categories

Types of NoSQL DB by Data Model

Data Model	Key Characteristics
Document	<ul style="list-style-type: none">• Store data in documents. Use structure like JSON• Intuitive way, closely aligned with OO programming (docs are objects)• Each record and its associated data are stored together, rather than distributed across tables.
Key-Value and Column-Family	<ul style="list-style-type: none">• Most basic types of non-tabular DB• Data stored as an attribute name (key), with its value• Value is entirely opaque, data can only be queried by the key.
Graph	<ul style="list-style-type: none">• Store data in graphs structures with nodes, edges and properties• Data modelled as a network of relationships between specific elements. (e.g. social networks)

Main Factor for No-SQL Emergence: Clusters



- Increase of scale (what is now called Big Data), produced a need of more computing resources
- Two options
 - **Scaling up (vertical)**: bigger machines, more processors, disk, storage and memory. More expensive! (also limits)
 - **Scaling out (horizontal)**: use a lots of small machines in a cluster. Much cheaper, and more resilient (keep going despite failures)
- RDB are not designed to run efficiently on clusters, while NoSQL have been designed to run on clusters

Distribution Models: Summary

- Replication

- Copies data across multiple servers
- Each data can be found in multiple places
- Synchronous or Asynchronous
- Asynchronous
 1. Primary-secondary
 2. Peer-to-peer

- Sharding

- Distributes different data across multiple servers.
- Each server acts as the single source for a subset of data


- Combining Replication with Sharding

Replication

- Replication provides redundancy and increases data availability.
- **Mainly** for resilience - in a big network, things WILL fail. Your system needs to continue as if nothing happened (Fault Tolerance)
- Also helps with performance
 - Reduce data distance travel by accessing local replica
 - Load balancing

Replication

ID	Name	Email
1	Bill	bill@mail.com
2	John	john@mail.com




ID	Name	Email
1	Bill	bill@mail.com
2	John	john@mail.com

ID	Name	Email
1	Bill	bill@mail.com
2	John	john@mail.com





Replication Resilience

ID	Name	Email
1	Bill	bill@mail.com
2	John	john@mail.com



ID	Name	Email
1	Bill	bill@mail.com
2	John	john@mail.com

ID	Name	Email
1	Bill	bill@mail.com
2	John	john@mail.com



Types of Replication

- Synchronous

- All replicas are updated on every write
- Reads are guaranteed to be up to date, so you can read from any node
- Maintains benefit of resilience, but can be too slow for some applications
- Used if reads MUST be up to date
- Works best if fewer writes are made

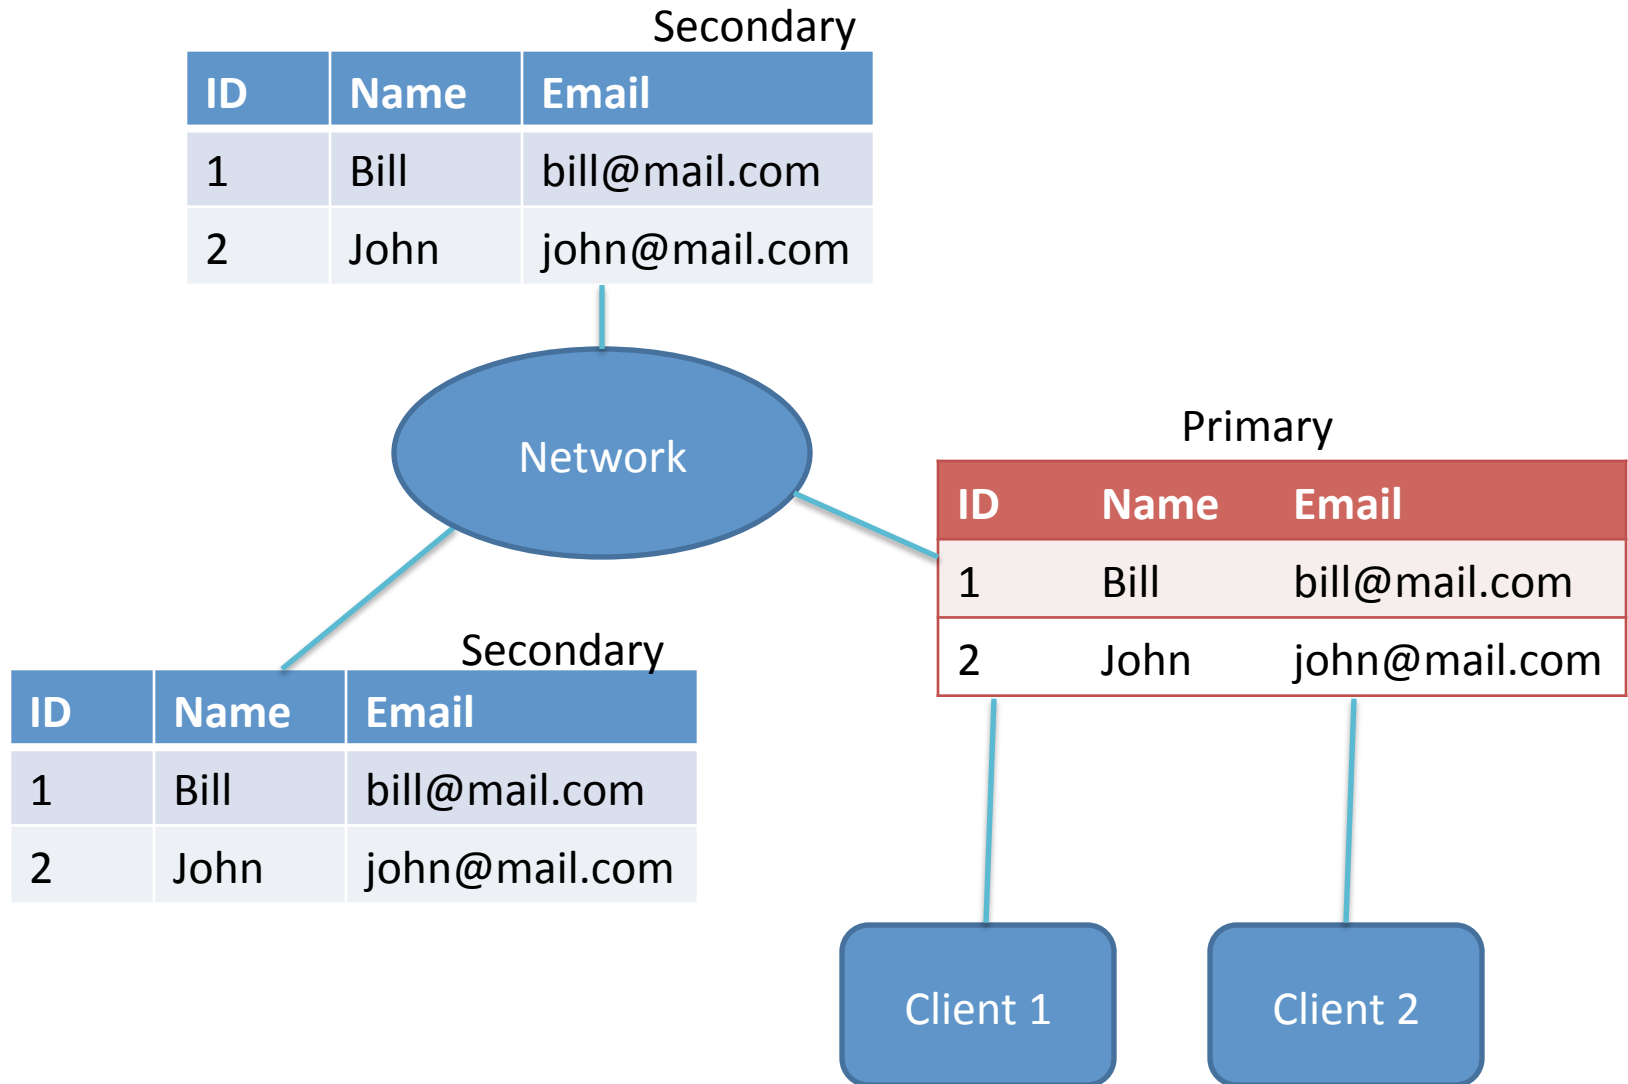
Types of Replication

- Asynchronous
 - Writes are propagated as soon as possible, but reads do not have to wait
 - Reads can be out of date
 - Eventual consistency
 - Works best if reads can be a little out of date
 - Methods include
 - Primary-secondary
 - Peer-to-peer

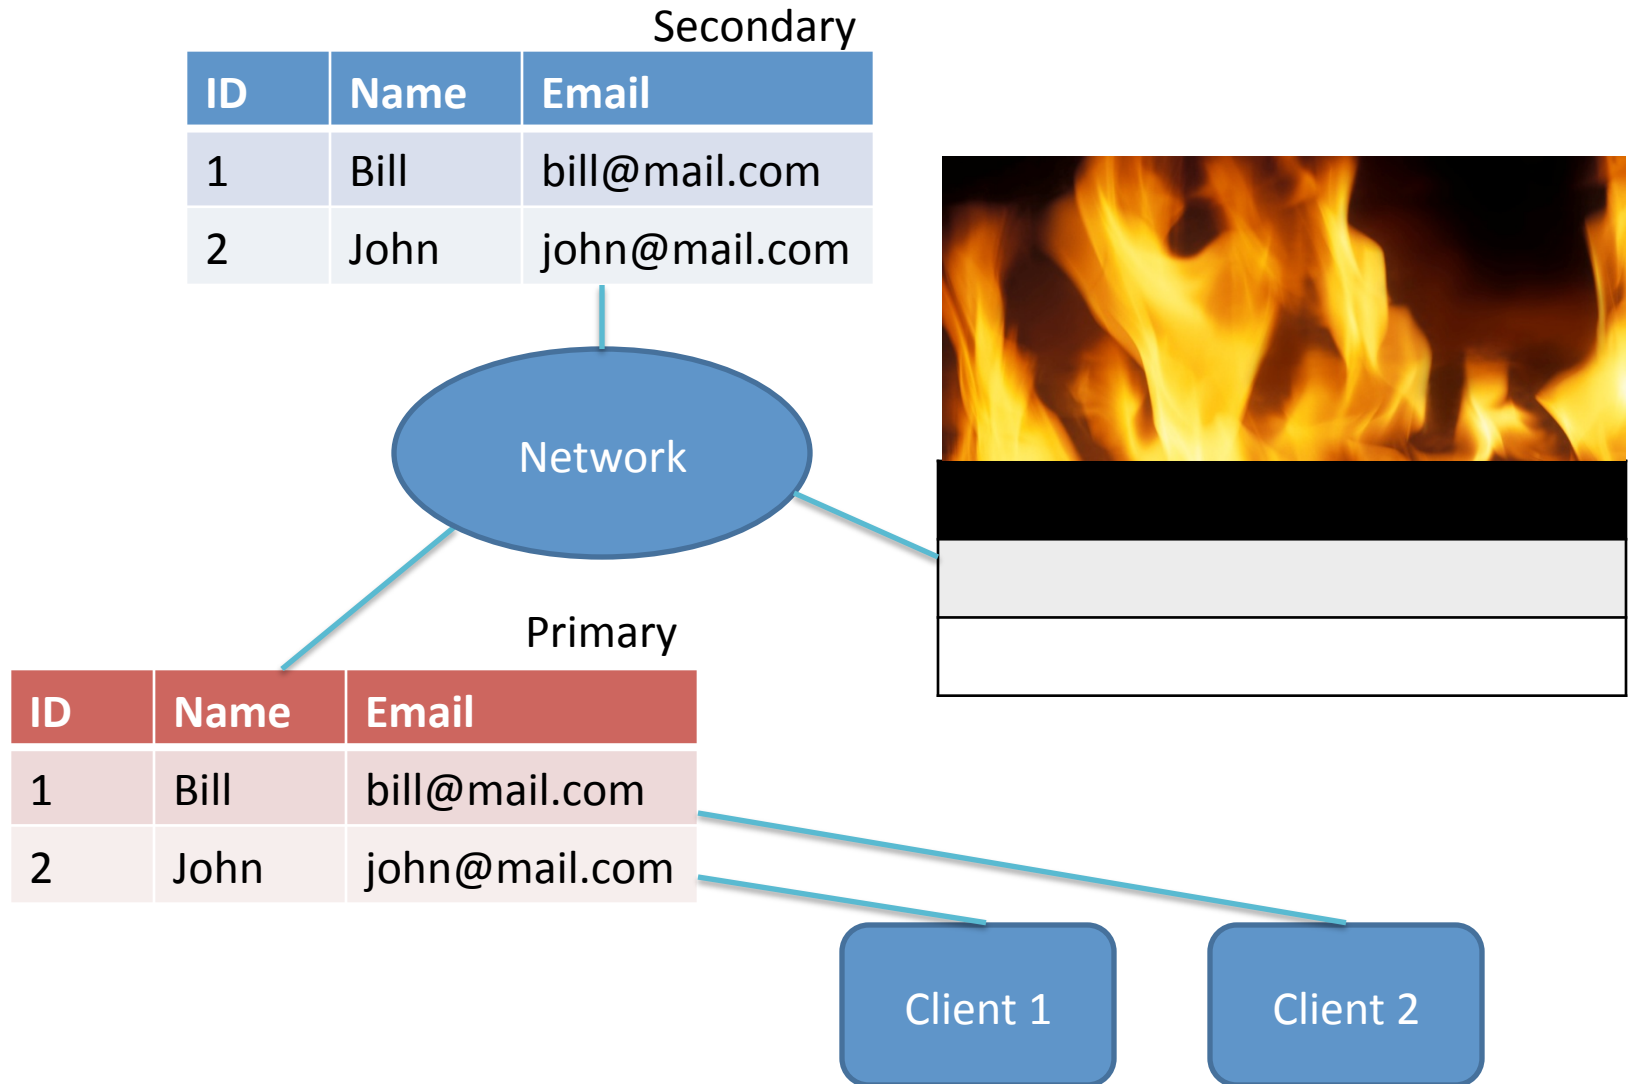
Primary-secondary Method

- A primary node, and 1 or more secondary nodes
- Not a SPOF (single point of failure) – if it fails, one of the secondaries will be promoted to primary
- Writes go only to the primary and are then propagated
- Secondaries can be read but not written
- Replication factor is the number of replicas made
- Generally, it is considered safe to have three copies in total

Primary – secondary replication



Primary in MongoDB



Consistency

- In a relational database, ACID consistency is handled by transactions, which maintain database integrity
- In NoSQL, consistency refers to whether or not reads reflect previous writes
 - **Strict consistency** – A read is guaranteed to get the most up to date data
 - **Eventual consistency** – Read data may be stale, but writes will catch up pretty quick

Back to MongoDB

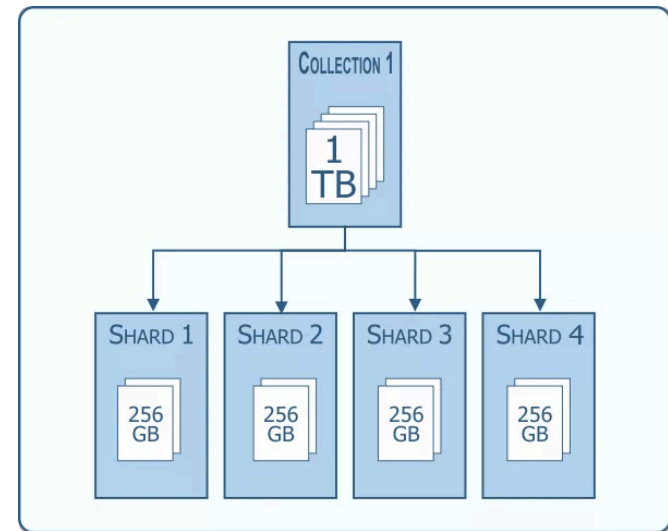
- Always reading from the primary gives strict consistency
- Reads from a secondary may be stale, so this is eventual consistency
- Why use eventual consistency?
 - For performance, reads can be spread across multiple secondary nodes
 - Some applications, e.g. Facebook don't need strict consistency
 - Offline analytics should read from secondaries to avoid overloading primary

Peer-to-Peer Replication

- The primary node can be a write bottleneck
- In systems with a high write rate, this can cause latency
- Peer-to-peer replication allows all nodes to accept reads and writes (there is no primary node)
- Fast and resilient
- Problems with inconsistency – particularly write inconsistency where two peers receive conflicting updates

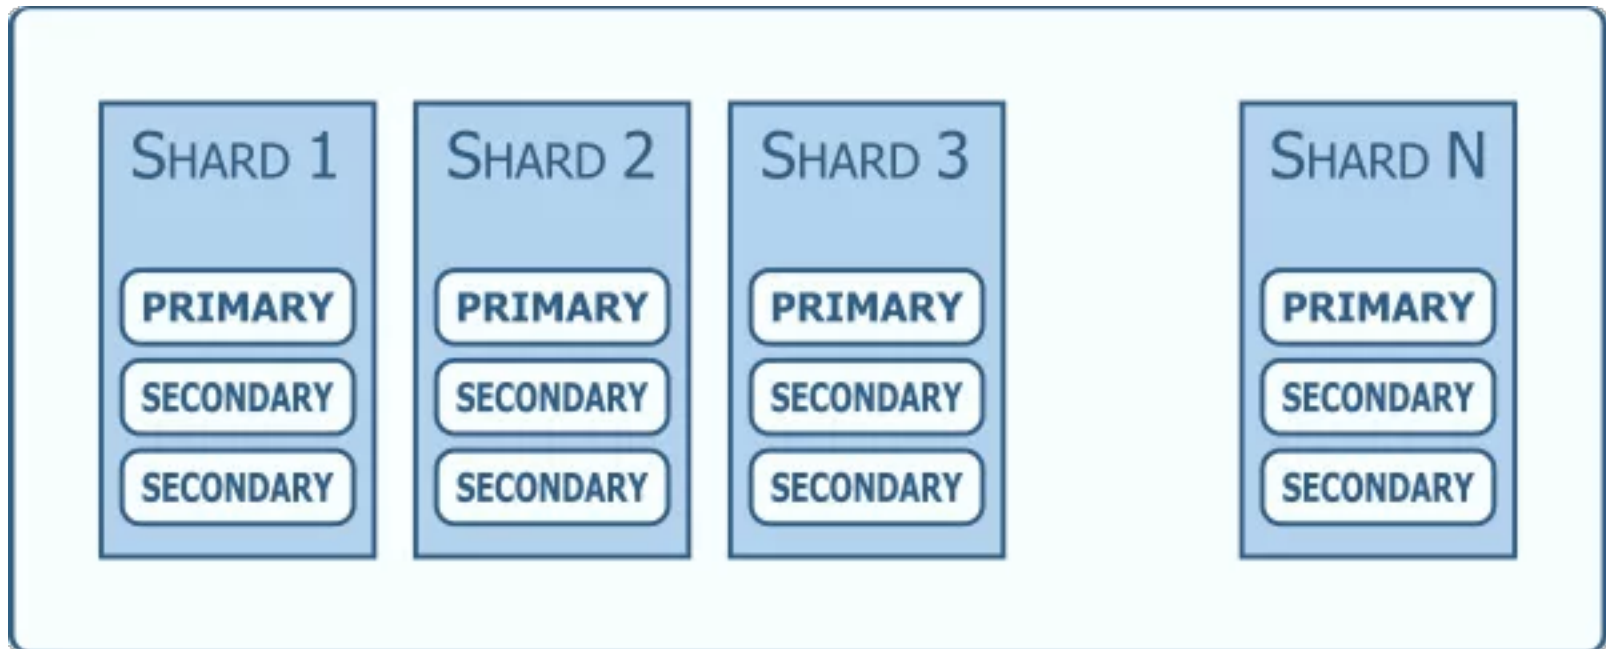
Sharding

- Spreads the data across nodes in the cluster
- Each node gets a subset of the data
- Might split the data by location (UK customer data in Ireland, American customer data in NY)
- Aggregate model helps sharding work



In MongoDB data is sharded at a collection level

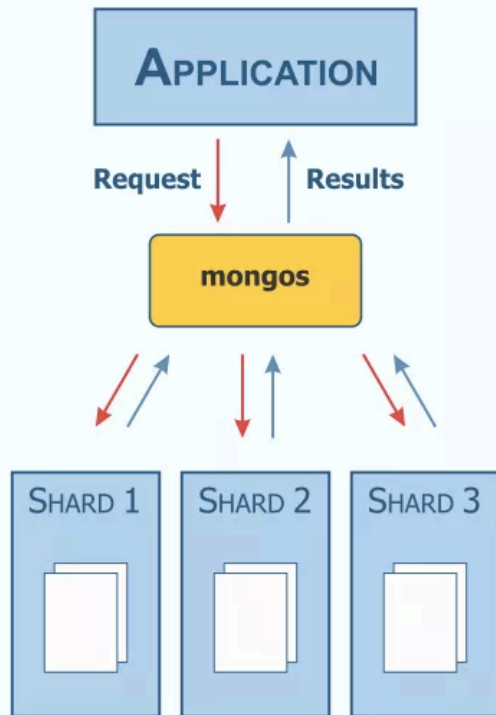
Combining Sharding with Relication



- A *replica set* is a group of mongod instances that maintain the same data set.
- Each cluster is, to keep redundancy, areplica sets with primary and secondary servers.

<https://www.kenwalger.com/blog/nosql/mongodb/mongodb-horizontal-scaling-sharding/>

Cluster configuration with MongoDB

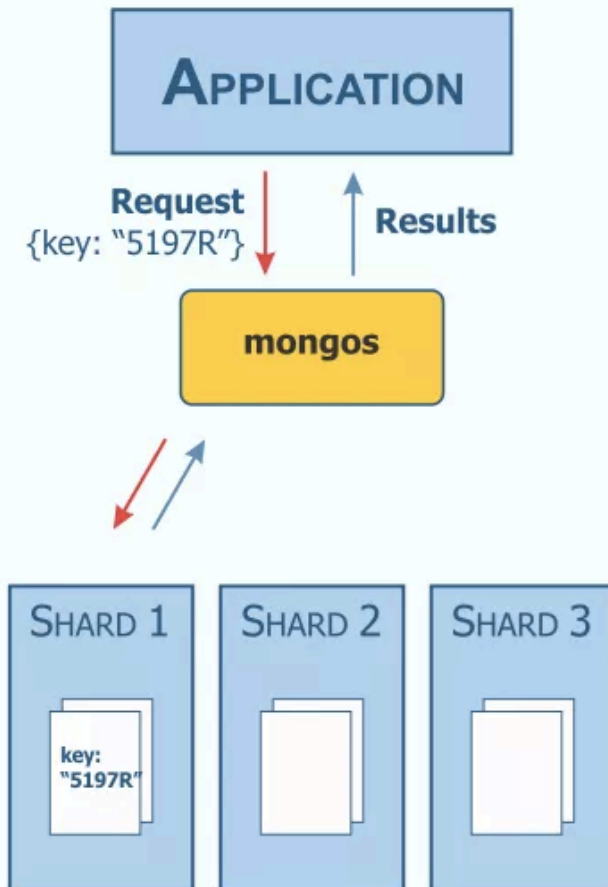


- **mongos** for “MongoDB Shard,” is a routing service for MongoDB shard configurations
- **mongos** acts as the interface between application and the data
- It routes queries and write operations to the appropriate shards
- An application access the data through **mongos** not touching the data directly
- Queries are directed to all shards unless it can be determined that the data resides on a particular shard

Shard Keys

- There has to be a better way than broadcasting that request to all shards
- **Use a Key**: to determine how documents in a collection are distributed across the shards.
- The Key is based on one or more chosen fields – much like a primary key in a RDB
- Shard keys can be chosen by hand but most systems offer auto-sharding (including resharding to maintain load balance)

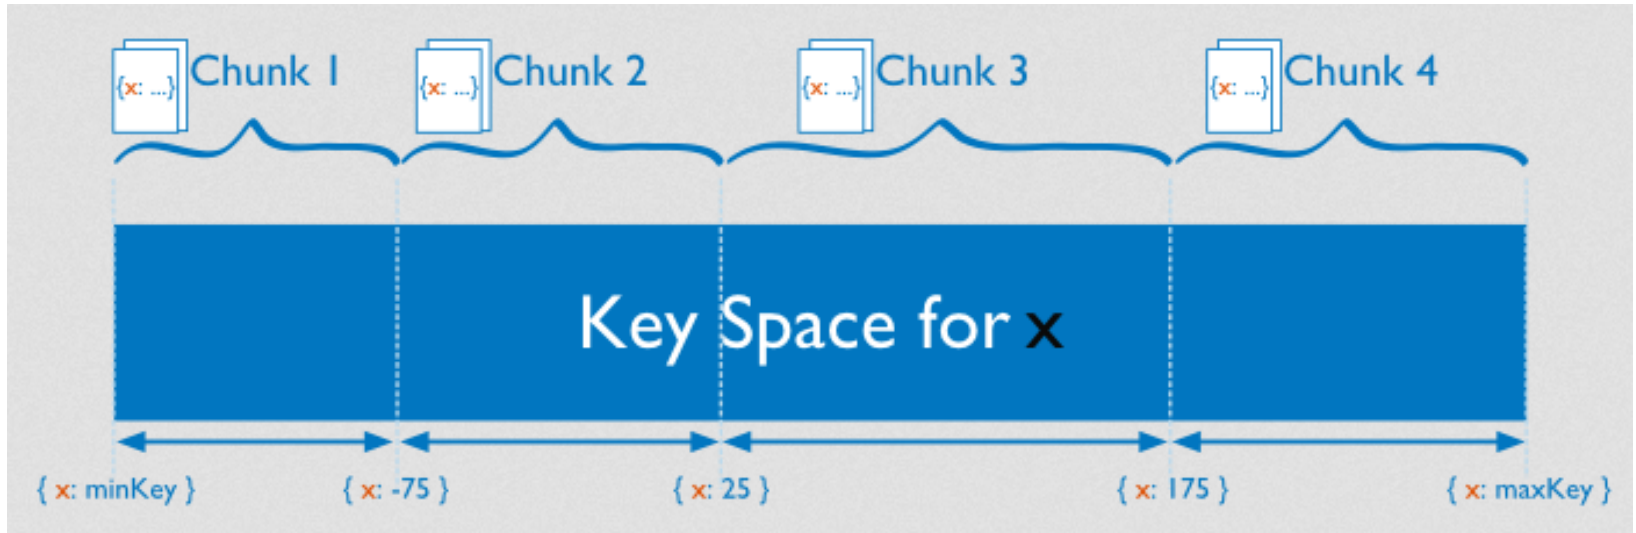
- If provided with a shard key during a query, the mongos knows how to route the request. More efficient!
- The choice of shard key affects the performance, efficiency, and scalability of a sharded cluster.



Choosing a Shard Key

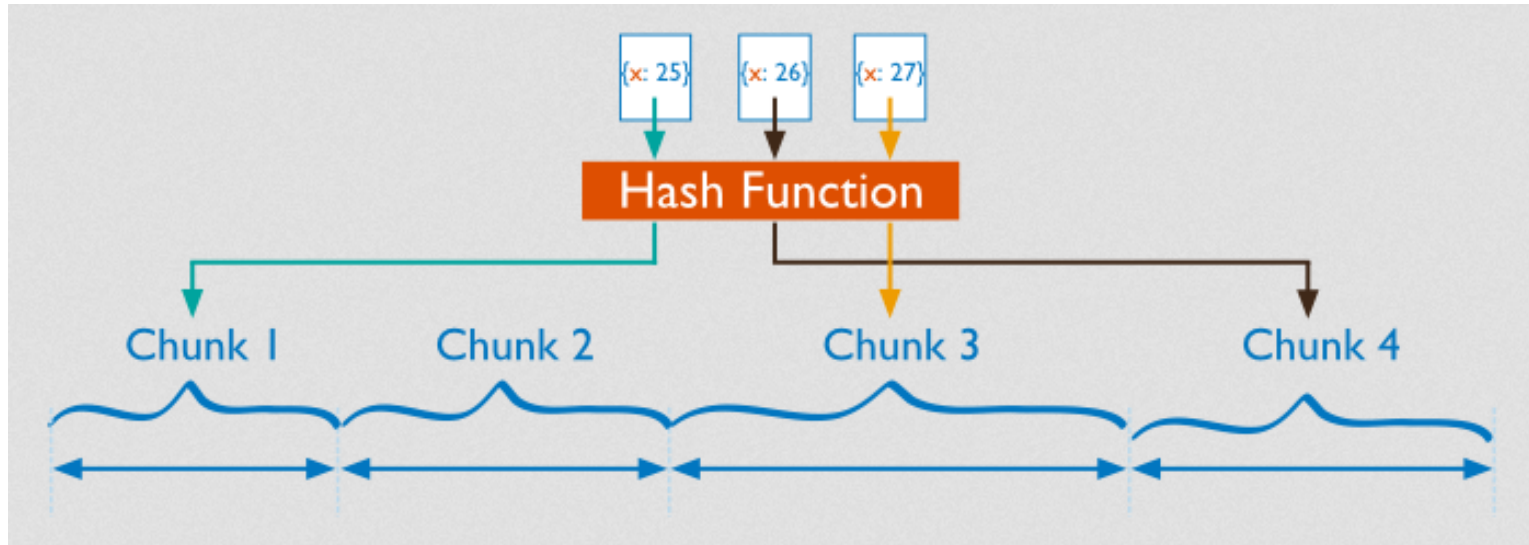
- Field(s) must appear in every record
- Easily divisible, with wide number of possible values
- Uniformly distributed across records and queries
- Heavily dependent on your application
- Think about fields that appear in common queries

Range Based Keys



- Divides the data set into ranges determined by the shard key values to provide **range based partitioning**.
- Consider a numeric shard key: If you visualize a line that goes from negative infinity to positive infinity, each value of the shard key falls at some point on that line.
- The line can be divided into smaller, non-overlapping ranges called **chunks** where a chunk is range of values from some minimum value to some maximum value.
- Documents with “close” shard key values are likely to be in the same chunk, and therefore on the same shard.

Hash Based Keys



- Compute a hash of a field's value, and then uses these hashes to create chunks (*hash based partitioning*).
- Two documents with “close” shard key values are *unlikely* to be part of the same chunk.
- This ensures a more random distribution of a collection in the cluster.

Summary

- Replication
 - Copies data across multiple servers
 - Each data can be found in multiple places
 - Synchronous or Asynchronous
 - Asynchronous
 1. Primary-secondary
 2. Peer-to-peer
- Sharding
 - Distributes different data across multiple servers.
 - Each server acts as the single source for a subset of data
- Combining Replication with Sharding (MongoDB)