

CSCU9YQ - NoSQL Databases

Lecture 3: The Aggregation Pipeline

Gabriela Ochoa



What is Data Aggregation?

- Any process in which information is gathered and expressed in a summary form, for purposes such as statistical analysis
- Purpose: get more information about particular groups based on specific variables such as age, profession, or income.
- Is a type of more sophisticated Query
- Aggregation operations
 - group values from multiple documents together
 - perform a variety of operations on the grouped data to return a single result.

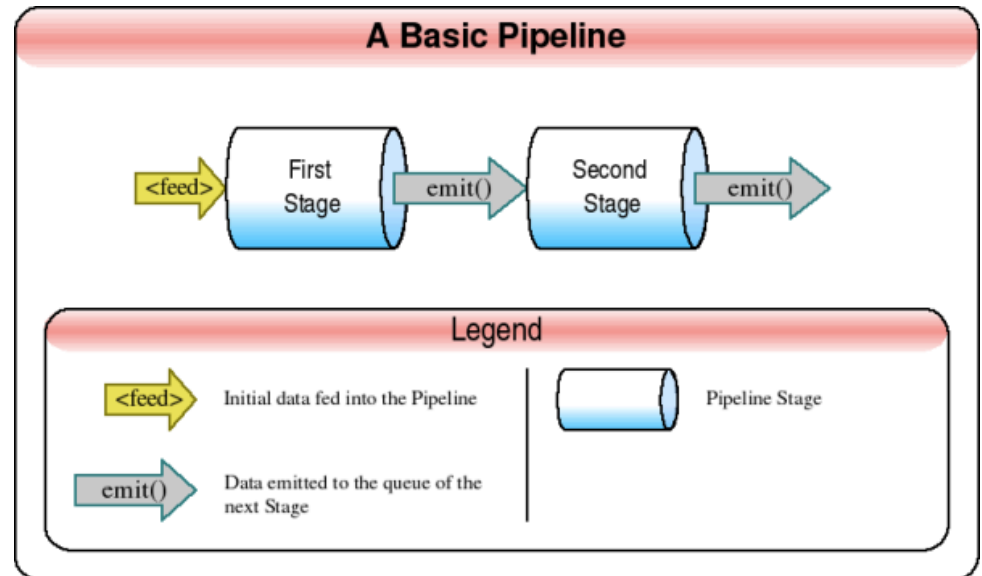
Aggregation in MongoDB

MongoDB provides three ways to perform aggregation

- **Aggregation pipeline**
 - Preferred method, use native code
 - More efficient than map-reduce
- **Map-reduce function**
 - Harder to program, can be more flexible
 - Less efficient, uses custom JavaScript functions
- **Single purpose aggregation methods**
 - Limited scope
 - Easy to program common aggregations

The Aggregation Pipeline

- Documents enter a multi-stage pipeline that transforms the documents into an aggregated result.
- The most basic pipeline stages provide
 - filters that operate like queries and
 - document transformations that modify the form of the output document.



Other pipelines operators provide tools for

- Grouping and sorting documents by specific field(s)
- Aggregating the content of arrays, including arrays of documents

Aggregation Pipeline



- Stages do not need to produce one output document for every input
- Some stages may generate new documents or filter out documents
- Mongo Shell method
 - `db.collection.aggregate([{ <stage> }, ...])`
 - Parameters in JSON format.
 - Stages appear in an array

Some Common Stages

Stage	Description
\$match:	Filter the collection according to the query parameters, and only pass through the documents matching the query, to the next stage of the pipeline.
\$group	Groups documents by some specified expression and outputs to the next stage a document for each distinct grouping.
\$project	Reshapes each document in the stream, such as by adding new fields or removing existing fields. For each input document, outputs one document.
\$unwind	Operates on arrays. Deconstructs an array field from the input documents to output a document for each element.
\$sort	Reorders the document stream by a specified sort key. Only the order changes; the documents remain unmodified.

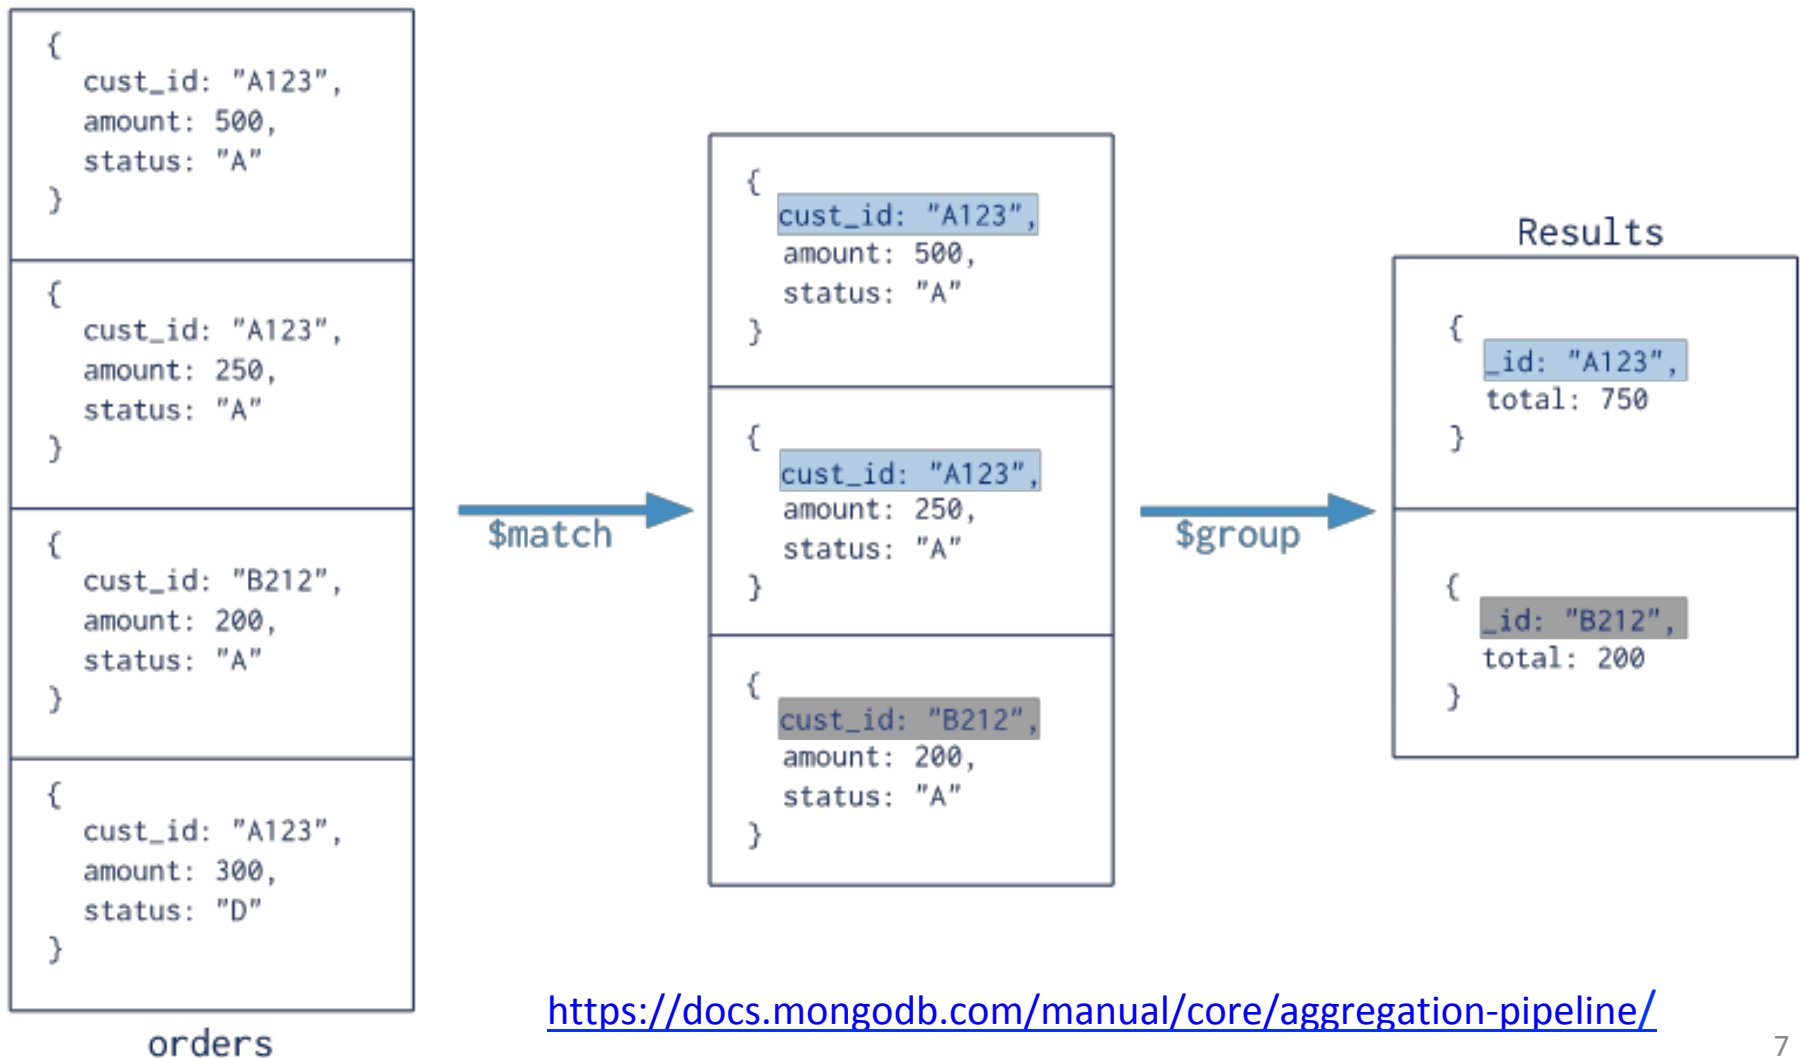
More stages:

<https://docs.mongodb.com/manual/reference/operator/aggregation-pipeline/#aggregation-pipeline-operator-reference>

Collection

↓

```
db.orders.aggregate( [
  $match stage → { $match: { status: "A" } },
  $group stage → { $group: { _id: "$cust_id", total: { $sum: "$amount" } } }
] )
```



\$match

- Filter documents
 - pass only the documents that match the specified condition(s) to the next pipeline stage.
- Should be placed early in the pipeline, as it limits the number of documents, and thus optimise the process
- Syntax (use existing query syntax): `{ $match: { <query> } }`

Form & Examples	Description
<code>{ <field1>: <value1>, ... }</code>	Equality condition
<code>{ <field1>: { <operator1>: <value1> }, ... }</code>	Conditions using query operators
<code>{ status: "A", qty: { \$lt: 30 } }</code>	Concatenation separated by ',' gives an implicit AND condition
<code>{ \$or: [{ status: "A" }, { qty: { \$lt: 30 } }] }</code>	OR condition needs operator \$or

\$match simple match


```
{  
  subject: "Hello There",  
  words: 218,  
  from: "norberto@mongodb.com"  
}
```

```
{  
  subject: "I love Hofbrauhaus",  
  words: 90,  
  from: "norberto@mongodb.com"  
}
```

```
{  
  subject: "MongoDB Rules!",  
  words: 100,  
  from: "hipster@somemail.com"  
}
```



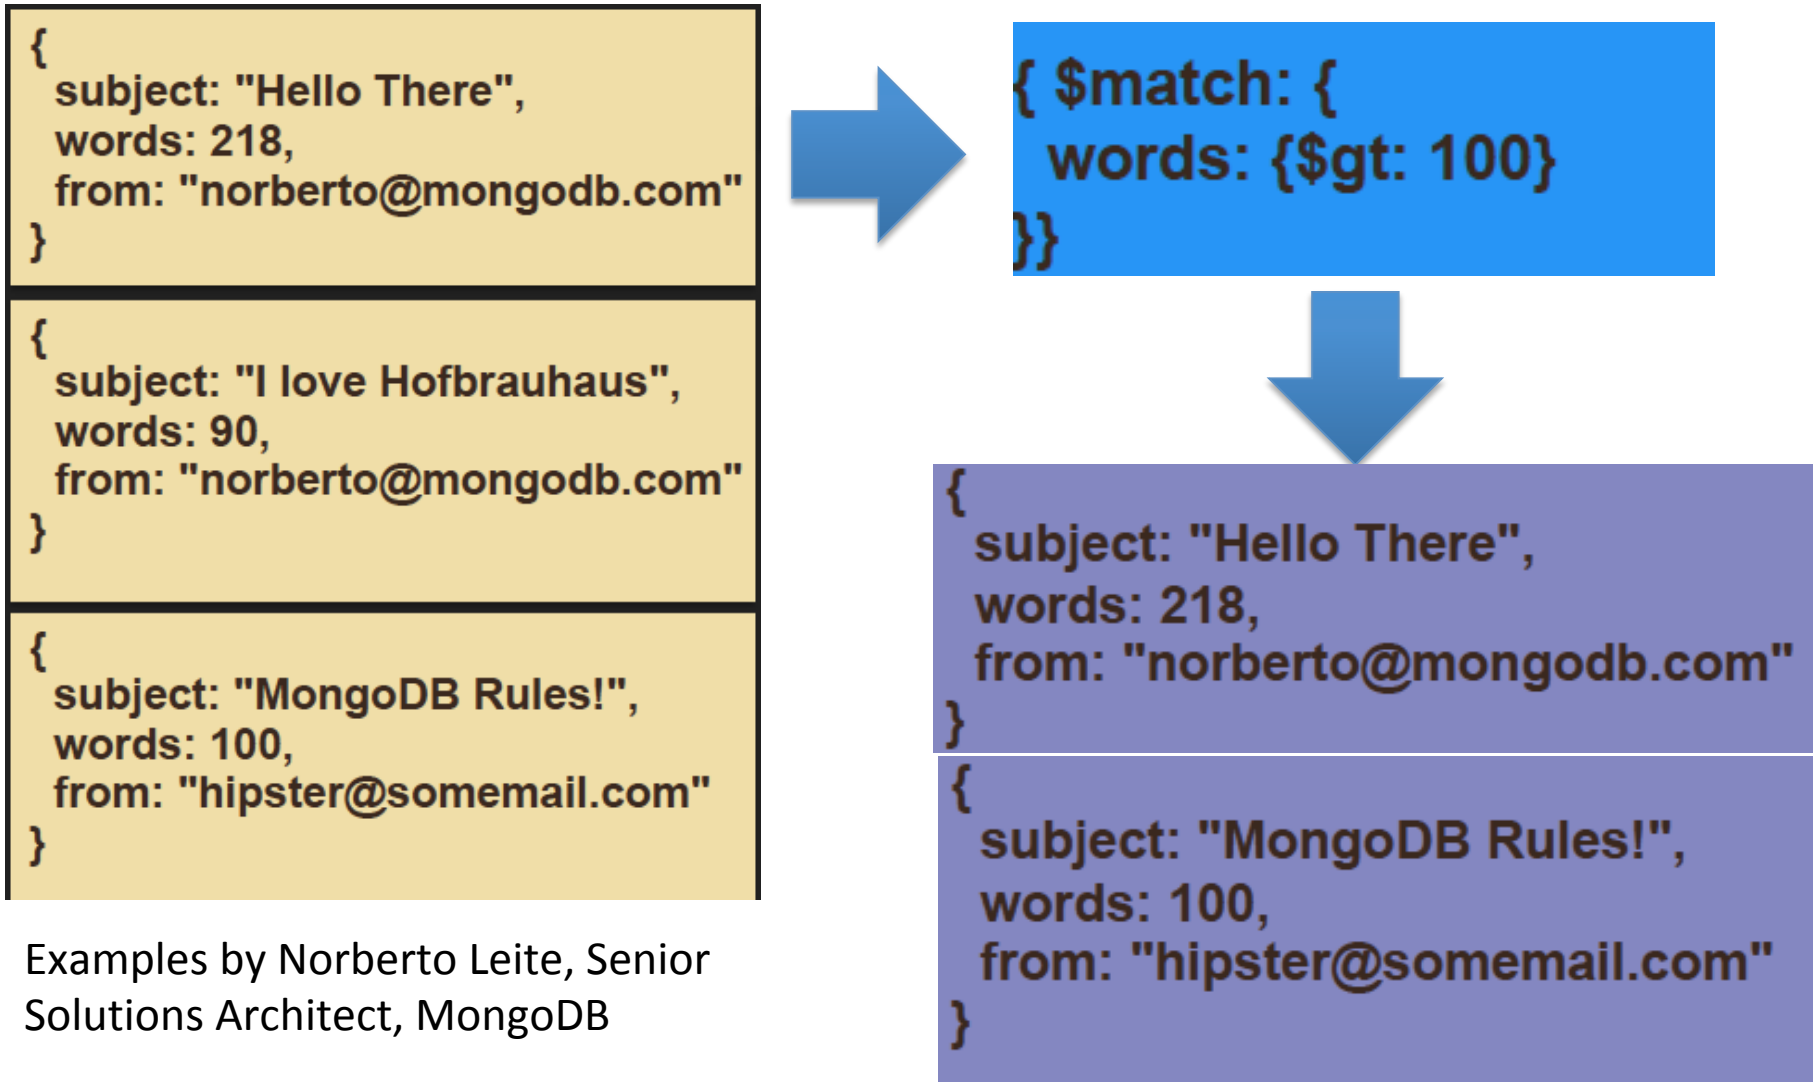
```
{ $match: {  
  from: "hipster@somemail.com"  
}}
```



```
{  
  subject: "MongoDB Rules!",  
  words: 100,  
  from: "hipster@somemail.com"  
}
```

Examples by Norberto Leite, Senior
Solutions Architect, MongoDB

\$match with Query Operations



\$project

- Reshape Documents
 - Include, exclude or rename fields
 - Inject computed fields
 - Create sub-document fields
- Syntax: { \$project: { <specification(s)> } }


Form	Description
<field>: <1 or true>	Specifies the inclusion of a field.
<field>: <0 or false>	Specifies the exclusion of a field.
_id: <0 or false>	Specifies the suppression of the _id field. By default the _id is included.
<field>: <expression>	Adds a new field or resets the value of an existing field.

\$project: Including and Excluding Fields

```
{  
  _id: 12345,  
  subject: "Hello There",  
  words: 218,  
  from: "norberto@mongodb.com"  
  to: [ "marc@mongodb.com",  
        "sam@mongodb.com" ],  
  account: "mongodb mail",  
  date: ISODate("2012-08-05"),  
  replies: 3,  
  folder: "Inbox",  
  ...  
}
```



```
{ $project: {  
  _id: 0,  
  subject: 1,  
  from: 1  
}}
```



```
{  
  subject: "Hello There",  
  from: "norberto@mongodb.com"  
}
```


Examples by Norberto Leite, Senior
Solutions Architect, MongoDB

\$project: Renaming and Computing Fields

```
{
  _id: 12345,
  subject: "Hello There",
  words: 218,
  from: "norberto@mongodb.com"
  to: [ "marc@mongodb.com",
        "sam@mongodb.com" ],
  account: "mongodb mail",
  date: ISODate("2012-08-05"),
  replies: 3,
  folder: "Inbox",
  ...
}
```



```
{ $project: {
  spamIndex: {
    $mul: ["$words", "$replies"]
  },
  user: "$from"
}}
```



```
{
  _id: 12345,
  spamIndex: 72.6666 ,
  user: "norberto@mongodb.com"
}
```

\$mul : operator multiply


Examples by Norberto Leite, Senior
Solutions Architect, MongoDB

\$project: Creating Sub-Document Fields

```
{
  _id: 12345,
  subject: "Hello There",
  words: 218,
  from: "norberto@mongodb.com"
  to: [ "marc@mongodb.com",
        "sam@mongodb.com" ],
  account: "mongodb mail",
  date: ISODate("2012-08-05"),
  replies: 3,
  folder: "Inbox",
  ...
}
```



```
{ $project: {
  subject: 1,
  stats: {
    replies: "$replies",
    from: "$from",
    date: "$date"
  }
}}
```



```
{
  _id: 375,
  subject: "Hello There",
  stats: {
    replies: 3,
    from: "norberto@mongodb.com",
    date: ISODate("2012-08-05")
  }
}
```

\$group

- Groups documents by some specified expression and outputs to the next state a document for each distinct grouping
- The output documents contain
 - an `_id` field with the distinct groups by key
 - Computed fields that hold the values of some accumulator expression grouped by the group's `_id` field.
 - `$max`, `$min`, `$avg`, `$sum`
 - `$addToSet`, `$push`
 - `$first`, `$last`

\$group syntax

```
{ $group: { _id: <expression>, <field1>: { <accumulator1> : <expression1> }, ... } }
```

- The `_id` field is mandatory; however, you can specify an `_id` value of null to calculate accumulated values for all the input documents as a whole.
- The remaining computed fields are optional and computed using the `<accumulator>` operators.
- The `_id` and the `<accumulator>` expressions can accept any valid expression.
- Valid Expressions
 - Field paths: "\$<field>" \$ followed by field name (or dotted field name)
 - Literals: of any type
 - Expression object: { <field1>: <expression1>, ... }

\$group: Calculating An Average

```
{  
  subject: "Hello There",  
  words: 218,  
  from: "norberto@mongodb.com"  
}
```

```
{  
  subject: "I love Hofbrauhaus",  
  words: 90,  
  from: "norberto@mongodb.com"  
}
```

```
{  
  subject: "MongoDB Rules!",  
  words: 100,  
  from: "hipster@somemail.com"  
}
```



```
{ $group: {  
  _id: "$from",  
  avgWords: { $avg: "$words" }  
}}
```



```
{  
  _id: "norberto@mongodb.com",  
  avgPages: 154  
}
```

```
{  
  _id: "hipster@somemail.com",  
  avgPages: 100  
}
```

Examples by Norberto Leite, Senior
Solutions Architect, MongoDB

\$group: Summing Fields and Counting

```
{  
  subject: "Hello There",  
  words: 218,  
  from: "norberto@mongodb.com"  
}
```

```
{  
  subject: "I love Hofbrauhaus",  
  words: 90,  
  from: "norberto@mongodb.com"  
}
```

```
{  
  subject: "MongoDB Rules!",  
  words: 100,  
  from: "hipster@somemail.com"  
}
```



```
{ $group: {  
  _id: "$from",  
  words: {$sum: "$words"},  
  mails: {$sum: 1}  
}}
```



```
{  
  _id: "norberto@mongodb.com",  
  words: 308,  
  mails: 2  
}
```

```
{  
  _id: "hipster@somemail.com",  
  words: 100,  
  mails: 1  
}
```

Examples by Norberto Leite, Senior
Solutions Architect, MongoDB

\$unwind


- Operates on an array field
- Deconstructs the array from the input documents
- Creates an output a document for each array element.
- Syntax: { \$unwind: <field path> }
 - <field path> is a string "\$<field_name>"

\$unwind: Collecting Distinct Values

```
{
  _id: 2222,
  subject: "2.8 will be great!",
  to: [ "marc@mongodb.com",
        "eliot@mongodb.com",
        "asya@mongodb.com",
      ],
  account: "mongodb mail"
}
```



```
{ $unwind: "$to" }
```



```
{ subject: "2.8 will be great!",
  to: "marc@mongodb.com",
  account : "mongodb mail" }
```

```
{ subject: "2.8 will be great!",
  to: "eliot@mongodb.com",
  account : "mongodb mail" }
```

```
{ subject: "2.8 will be great!",
  to: "asya@mongodb.com",
  account : "mongodb mail" }
```

\$sort

- Sorts all input documents and returns them to the pipeline in sorted
- Syntax:
 - { \$sort: { <field1>: <sort order>, <field2>: <sort order> ... } }
 - The sort order is a value of 1 or -1 to specify an ascending or descending sort respectively

Example combining \$project and \$sort: Sports Club

```
{
  _id : "jane",
  joined : ISODate("2011-02-02"),
  likes : ["golf", "racquetball"]
}
```

```
{
  _id : "joe",
  joined : ISODate("2012-07-02"),
  likes : ["tennis", "golf", "swimming"]
}
```

user collections (more documents)

...

Returns user names in upper case and in alphabetical order

Example from : MongoDB Documentation
<https://docs.mongodb.com/manual/tutorial/aggregation-with-user-preference-data/>



```
[
  { $project : { name:{ $toUpper:"$_id" }, _id:0 } },
  { $sort : { name : 1 } }
]
```



```
{
  "name" : "JANE"
},
```

```
{
  "name" : "JILL"
},
```

```
{
  "name" : "JOE"
}
```

Example combining \$project and \$sort: Sports Club

```
{
  _id : "jane",
  joined : ISODate("2011-02-02"),
  likes : ["golf", "racquetball"]
}
```

```
{
  _id : "joe",
  joined : ISODate("2012-07-02"),
  likes : ["tennis", "golf", "swimming"]
}
```

user collections (more documents)

...

- Returns user names sorted by the month they joined.
- \$month operator converts the values of the joined field to integer representations of the month.



```
[
  { $project :
    {
      month_joined : { $month : "$joined" },
      name : "$_id",
      _id : 0
    }
  },
  { $sort : { month_joined : 1 } }
]
```



```
{
  "month_joined" : 1,
  "name" : "ruth"
},
```

```
{
  "month_joined" : 2,
  "name" : "jane"
},
```

Single Purpose Aggregation Operations

- MongoDB also provides operations for common aggregations
 - `db.collection.count(query, options)`: Returns the count of documents that would match a `find()` query for the collection or view.
 - `db.collection.estimatedDocumentCount(options)`: returns the count of all documents in a collection. Wraps the count command.
 - `db.collection.distinct(field, query, options)`
- Aggregate documents from a single collection
- Provide simple access to common aggregation processes, they lack the flexibility and capabilities of the aggregation pipeline and map-reduce.

Collection



```
db.orders.distinct( "cust_id" )
```

<pre>{ cust_id: "A123", amount: 500, status: "A" }</pre>
<pre>{ cust_id: "A123", amount: 250, status: "A" }</pre>
<pre>{ cust_id: "B212", amount: 200, status: "A" }</pre>
<pre>{ cust_id: "A123", amount: 300, status: "D" }</pre>

orders

distinct → ["A123", "B212"]

`db.collection.distinct(field, query, options)`

- field: string with the field for which to return distinct values
- query: document with a query that specifies the documents from which to retrieve the distinct values
- Options: a document that specify options

<https://docs.mongodb.com/manual/aggregation/>

Aggregation Summary

- Powerful tool to process Data, rich library of functions
- A series of Document Transformation, concatenated stages
- MongoDB provides three ways to perform aggregation
 - Aggregation pipeline
 - Preferred method, use native code
 - More efficient than map-reduce
 - Single purpose aggregation methods
 - Limited scope
 - Easy to program common aggregations
 - Map-reduce function (Next week)
 - Harder to program, can be more flexible
 - Less efficient, uses custom JavaScript functions