

CSCU9YQ - NoSQL Databases

Lecture 8: Cursors and Projections in MongoDB

Gabriela Ochoa

Cursors

- The `db.collection.find()` method returns a cursor.
- A cursor is a 'pointer' to the results.
- To access the documents, you need to iterate the cursor.

Cursors

- If the returned cursor is **not** assigned to a variable using the **var** keyword, then the cursor is automatically iterated up to 20 times (print up to the first 20 documents in the result)
- When assigned to a variable, to access the whole set of results, you need to
 - Iterate the cursor manually
 - Use the iterator index.

Manually Iterate the Cursor

```
var myCursor = db.users.find( { type: 2 } );  
  
myCursor
```

```
var myCursor = db.users.find( { type: 2 } );  
  
while (myCursor.hasNext()) {  
    print(tojson(myCursor.next()));  
}
```

```
var myCursor = db.users.find( { type: 2 } );  
  
while (myCursor.hasNext()) {  
    printjson(myCursor.next());  
}
```

- If you call the cursor variable in the shell, it will iterate up to 20 times and print the matching
- The shell will say: type 'it' for more
- Use the cursor method `next()` to access all the documents
- Print results with `tojson()`
- As an alternative print operation, use `printjson()` helper method to replace `print(tojson())`

Iterator Index

```
var myCursor = db.inventory.find( { type: 2 } );  
var documentArray = myCursor.toArray();  
var myDocument = documentArray[3];
```

- You can use the `toArray()` method to iterate the cursor and return the documents in an array
- The `toArray()` method
 - Loads into RAM all documents returned by the cursor
 - Exhausts the cursor.

Projections

- By default, queries in MongoDB return all fields in matching documents.
- To limit the amount of data, you can include a projection document to specify or restrict fields to return.

<https://docs.mongodb.com/manual/tutorial/project-fields-from-query-results/>

Example: Inventory data

```
db.inventory.insertMany( [  
  { item: "journal", status: "A", size: { h: 14, w: 21, uom: "cm" },  
    instock: [ { warehouse: "A", qty: 5 } ] },  
  { item: "notebook", status: "A", size: { h: 8.5, w: 11, uom: "in" },  
    instock: [ { warehouse: "C", qty: 5 } ] },  
  { item: "paper", status: "D", size: { h: 8.5, w: 11, uom: "in" },  
    instock: [ { warehouse: "A", qty: 60 } ] },  
  { item: "planner", status: "D", size: { h: 22.85, w: 30, uom: "cm" },  
    instock: [ { warehouse: "A", qty: 40 } ] },  
  { item: "postcard", status: "A", size: { h: 10, w: 15.25, uom: "cm" },  
    instock: [ { warehouse: "B", qty: 15 }, { warehouse: "C", qty: 35 } ] }  
]);
```

- If you do not specify a projection document, the `db.collection.find()` method returns all fields in the matching documents.
- `db.inventory.find({ status: "A" })`
 - Returns all fields from all documents where the status equals "A":

Return the Specified Fields

- A projection can explicitly include several fields by setting the <field> to 1 in the projection document.
- The `_id` field is shown by default
- `db.inventory.find({ status: "A" }, { item: 1, status: 1 })`

```
{ "_id": ObjectId("6shssgahs..."), item: "journal", status: "A" }  
{ "_id": ObjectId("3sfhf7jnsh..."), item: "notebook", status: "A" },  
{ "_id": ObjectId("8sasgahyv..."), item: "postcard", status: "A" }
```

- You can suppress the `_id` field by setting it to 0 in the projection
- `db.inventory.find({ status: "A" }, { item: 1, status: 1, _id: 0 })`

```
{ item: "journal", status: "A" }  
{ item: "notebook", status: "A" }  
{ item: "postcard", status: "A" }
```


Return all but the Excluded Fields

- You can use a projection to exclude specific fields
- Notice that, with the exception of the `_id` field, you **cannot** combine inclusion and exclusion statements in projection documents.
- `db.inventory.find({ status: "A" }, { status: 0, instock: 0 })`

```
{"_id": ObjectId("6shssgahs..."), item: "journal", size: { h: 14, w: 21, uom: "cm" }},  
{"_id": ObjectId("3sfhf7jnsh..."), item: "notebook", size: { h: 8.5, w: 11, uom: "in" }},  
{"_id": ObjectId("8sasgahyv..."), item: "postcard", size: { h: 10, w: 15.25, uom: "cm" }}
```

Return Specific Fields in Embedded Documents

- You can return specific fields in an embedded document.
- Use the dot notation to refer to the embedded field and set to 1 in the projection document.

```
db.inventory.find( { status: "A" }, { item: 1, status: 1, "size.uom": 1, _id: 0 } )
```

```
{item: "journal", status: : "A", size: { uom: "cm" }}  
{item: "notebook", status: "A", size: { uom: "cm" }}  
{item: "postcard", status: "A", size: { uom: "cm" }}
```

The uom field remains embedded in the size document.

Suppress Specific Fields in Embedded Documents

- You can return specific fields in an embedded document.
- Use the dot notation to refer to the embedded field and set to 0 in the projection document.

```
db.inventory.find( { status: "A" }, { "size.uom": 0, _id: 0 })
```

```
{item: "journal", status: "A", size: { h: 14, w: 21 },  
 instock: [ { warehouse: "A", qty: 5 } ] },  
{item: "notebook", status: "A", size: { h: 8.5, w: 11},  
 instock: [ { warehouse: "C", qty: 5 } ] },  
{ item: "postcard", status: "A", size: { h: 10, w: 15.25}  
 instock: [ { warehouse: "B", qty: 15 }, { warehouse: "C", qty: 35 }
```

Projection on Embedded Documents in an Array

- Use dot notation to project specific fields inside documents embedded in an array (use " ")

```
db.inventory.find( { status: "A" }, {item: 1, status: 1, "instock.qty": 1, _id: 0 })
```

```
{item: "journal", status: "A", instock: [{ warehouse: "A", qty: 5 } ] }  
{item: "notebook", status: "A", instock: [ { qty: 5 } ] }  
{ item: "postcard", status: "A", instock: [{qty: 15 }, {qty: 35 } ]}
```

Project Specific Array Elements in the Returned Array

- MongoDB provides the following projection operators for manipulating arrays:
 - **\$elemMatch**: limits the contents of an <array> field from the query results to contain only the first element matching the \$elemMatch condition.
 - **\$slice**: controls the number of items of an array that a query returns.
 - **\$**: limits the contents of an <array> from the query results to contain only the first element matching the query document.
- You **cannot** project specific array elements using the array index; e.g. { "instock.0": 1 }

Project Specific Array Elements in the Returned Array

- Example using the `$slice` projection operator to return the last element in the `instock` array

```
db.inventory.find( { status: "A" }, { _id:0, item: 1, status: 1, instock: { $slice: -1 } })
```

```
{item: "journal", status: "A", instock: [{ qty: 5 } ] }  
{item: "notebook", status: "A", instock: { warehouse: "C", qty: 5 } } ,  
{item: "postcard", status: "A", instock: [{ warehouse: "C", qty: 35 } ] }
```

Summary

- Cursors

- The `db.collection.find()` method returns a cursor.
- To access the whole set of results, you need to
 - Iterate the cursor manually
 - Use the iterator index.

- Projections

- Specify or restrict fields to return from a query
- Return or suppress specified fields
- Return or suppress specific fields in embedded documents
- Project embedded documents in an array
- Project specific array elements in the returned array