

# CSCU9YQ - NoSQL Databases

## Lecture 8: MongoDB Use Cases

Gabriela Ochoa

# NoSQL Use Cases

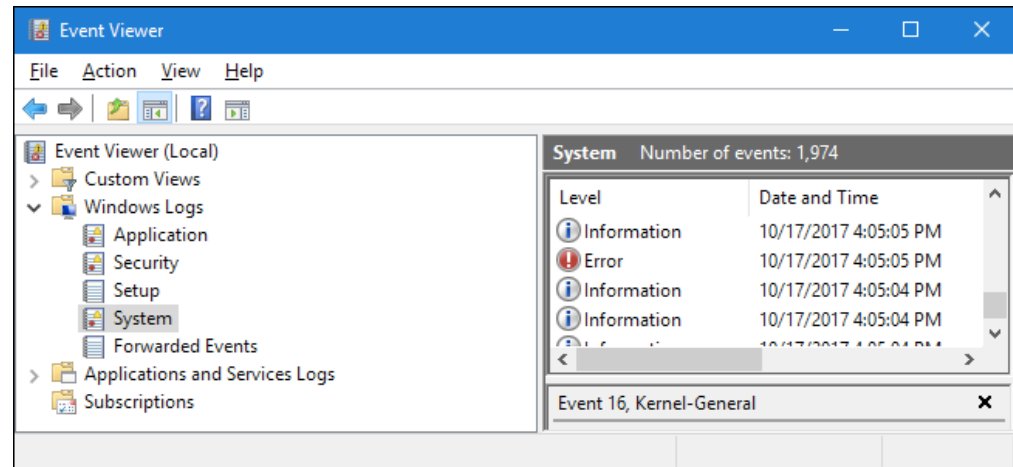
- Relational DB
  - will still be around for a long time
  - right kind of DB for handling centralised applications that require sophisticated transaction handling (e.g., general accounting).
- NoSQL DB
  - will continue growing in importance
  - better suited to support widely distributed cloud applications and their specific use cases

# Document DB Suitable Use Cases

- Event Logging
- Content Management Systems, Blogging Platforms
- Web Analytics or Real-Time Analytics
- E-commerce Applications

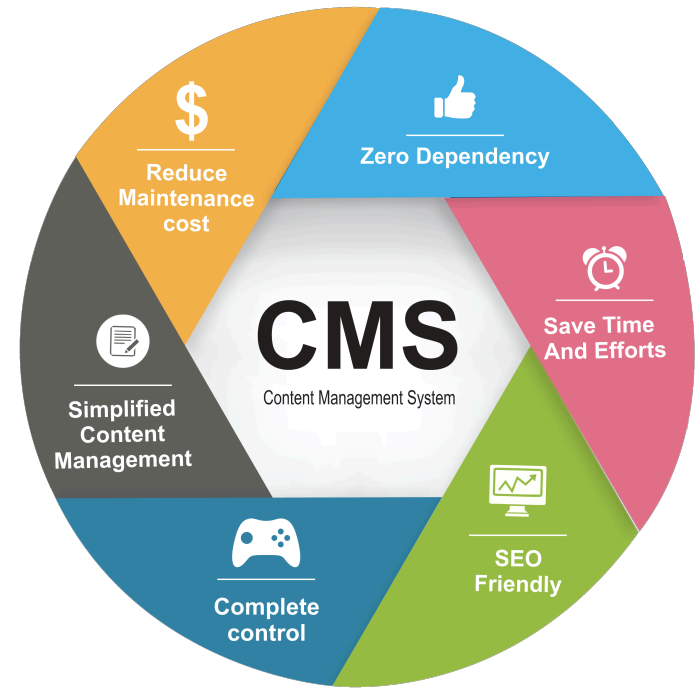
# Event Logging

- Different applications want to log events and errors each with their own format
- With a Document DB
  - Documents can store all these different types of events
  - Specially true when the data captured by the events keep changing (flexible schema)
  - Events can be sharded by the application that originated it, or by type of event (login, rt



# Content Management Systems

- Manage the creation and modification of digital content.
- Typically support multiple users in a collaborative environment
- With a Documents DB
  - JSON formant and flexible schema are useful
  - Easy to store user comments, registrations, profiles and other documents



# Web/Real Time Analytics

- Measurement, collection, analysis and reporting of web data for purposes of understanding and optimising web usage.
- With a Documents DB
  - Easy to store page views or unique visitors, as part of the document can be updated
  - New metrics can be easily added without schema changes.



# E-Commerce Applications

- Manage the creation and modification of digital content.
- Typically support multiple users in a collaborative environment
- With a Documents DB
  - JSON formant and flexible schema are useful
  - Easy to store user comments, registrations, profiles and other documents



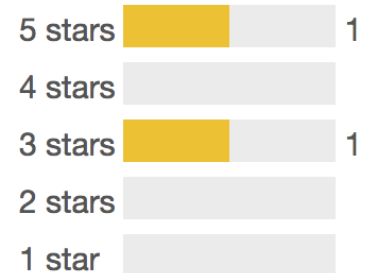
# Specific Example: Product Review Site

- Undergraduate project supervised by Dr. K. Swingler
- Web front end and Python, MongoDB at server
- Allows users to review products
- Instead of writing sentences, they are invited to enter key, value pairs that describe features of the product

## Reviews

**Average rating:** 4.00 out of 5 stars

**2 reviews**





# Product Review Site

- Example of data provided:
  - **Product**: JP Excite Ride Surfboard
  - **Value**: A little expensive, but worth it, 3/5
  - **Control**: Very easy to control, 5/5
  - **Speed**: Quite fast for a freestyle board, 4/5
- The identity of the products and the features that they contain cannot be known at design time, so **schemaless** design is needed

# JSON Object

```
{Product: "JP Excite Ride Surfboard",  
Category:"Windsurf boards",  
ReviewerID:2354242,  
{Value: {Description:"A little expensive,  
but worth it", Score:3/5}},  
{Control:{Description: "Very easy to  
control", Value:5/5}},  
{Speed: {Description:"Quite fast for a  
freestyle board", 4/5}}}
```

# Use Cases

- Defining use cases is essential in NoSQL design as it informs decisions of:
  - What the aggregate model should be
  - What fields are indexed
  - How data should be sharded

# Use Case 1

- *Description*: A user wants all the reviews of a single known product
- *Prevalence*: A common search
- *Suitable Aggregation*: Obvious document structure is one document per product with an array of review documents within it

# Use Case 2

- *Description*: A user knows the category of the product, but needs help choosing
- *Prevalence*: A common search
- *Suitable Aggregation*: Keeping all products in a category in one document is not practical
- *Suitable Sharding*: Keeping all products in a category on the same shard is sensible

# Use Case 3

- *Description*: A user wants to see all the reviews they (or another chosen user) have made
- *Prevalence*: A rare search
- *Suitable Aggregation*: To optimise this search, aggregate by user with all a user's reviews in the same document

# Decision

- Aggregate by product seems most sensible:

```
{Product: "JP Excite Ride Surfboard",  
Category: "Windsurf boards",  
Reviews: [. . .]}
```

# Embed or Link?

- Should the reviews array contain full reviews or links?
- To get all reviews in one read, best to embed
- No need for normalised form, as reviews rarely get updated and are stored only once
- Link to reviewer data, don't embed



# Solution





























```
{Product: "JP Excite Ride Surfboard",  
Category: "Windsurf boards",  
OverallScore: 4/5,  
Reviews:  
  [{ReviewerID: 2352,  
    Value: {Description: "Good", Score: 4/5},  
    Speed: {Description: "Slow", Score: 2/5}},  
    {ReviewerID: 2352,  
    Value: {Description: "Great", Score: 4/5},  
    Speed: {Description: "Fast", Score: 4/5}}  
  ]}
```

# Challenges: Knowing the keys

- If each review can have different field names, how can I know what to search for?
- Could maintain an array of keys in the product document:

```
{Product: "JP Excite Ride Surfboard",  
Category: "Windsurf boards",  
Keys: ["Value", "Speed", "Control"], ...}
```
- Similarly, when new reviews are entered, possible available keys can be shown to try and avoid synonyms

# Examples from MongoDB.org

Single View	Internet of Things	Mobile	Real-Time Analytics
   	   	   	   
Catalog	Personalization	Content Management	
   	   	   	

**Figure 1:** MongoDB Use Cases

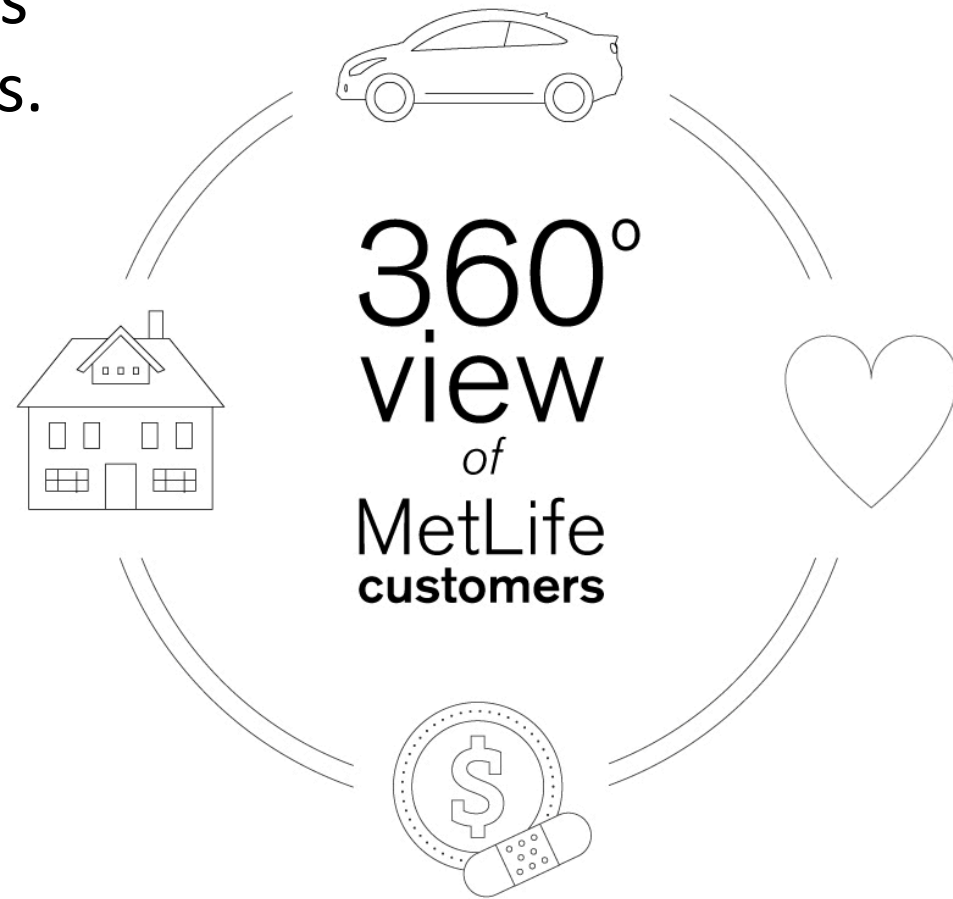
# Examples from MongoDB.org

- MongoDB White Papers
  - <https://www.mongodb.com/white-papers>
- Fortune 500 (world's 500 largest companies) companies and startups are using MongoDB
  - to create new types of applications
  - improve customer experience
  - accelerate time to market and reduce costs.

# MongoDB usage statistics

- Has been downloaded over 15 million times
- 2,000+ customers
- Significant customers
  - ADP, Bosch, Cisco, eBay, Expedia, MetLife, Telefonica, Ticketmaster and Verizon.
- Among the Fortune 500
  - Financial Services (40), Retailers (15), Telecomm (15), Technology (15), Healthcare (15), Electronics (10), Media Entertainment (10)

- MetLife is one of the world's largest insurance companies.
- Customer data was 'siloed' (isolated), difficult for call centre representatives to resolve customer issues efficiently
- 360 View developed with MongoDB in 3 months
- Mongo DB allowed the aggregation of data in a single data hub



**The 360-degree customer view:** creating a complete view by aggregating data from the various touch points that a customer may use to contact a company to purchase products and receive service and support.

Life (3) ▾

||| xxxxxxxx5456  
800-638-5433

||| xxxxxxxx2218  
800-638-5433

||| xxxxxxxx3476  
800-638-5433

TCA (1) ▾

||| xxxxxxxx9462  
800-638-7283

Confidence Level

||| High

||| Medium

||| Low



Policy #

xxxxxxx5456

Customer Service Phone Number

800-638-5433

Lucy Merryweather

Owner

SSN  
xxx-xx-xxxx

DOB  
August 10, 1973

Address  
109 West 93rd Avenue  
Lincoln, OK 07882

Phone  
(819) 555-7702

Alternate Phone  
(819) 555-2231

Email  
lkmerry@gmail.com

Product Type  
Life

Contract Type  
Whole Life

Group Name  
[No Data]

Sales Agent  
Tommy Topseller

Contract Status  
Active

Franchise Name  
New England Financial

Group Number  
[No Data]

Agent ID  
99B 560

Close Transaction Details

Transaction Details

Last 60 days • 2 transactions

All Transactions ▾

All Service Channels ▾

All Sources ▾

Self Service Address Change

Received: January 20, 2012 @ 10:00 am

Status: Completed | Source: Inbound | Channel: eService | System: BOSSLA

View Documents (2)

Change Confirmation (Mailed) • January 24, 2012 @ 4:31 pm

Change Request • January 20, 2012 @ 9:50 am

Self Service Bene Change

Received: June 1, 2011 @ 11:30 am

Status: Completed | Source: Inbound | Channel: eService | System: BOSSLA

View Documents (2)

Change Confirmation (Mailed) • June 3, 2011 @ 6:03 pm

Change Request • June 1, 2011 @ 11:20 am

The Wall App aggregates over 100 million customers, 100 products and over 70 source systems into a single data hub.

It presents the data in an intuitive, Facebook-like interface for customer service representatives

Viewed by MetLife and the industry as an overwhelming success

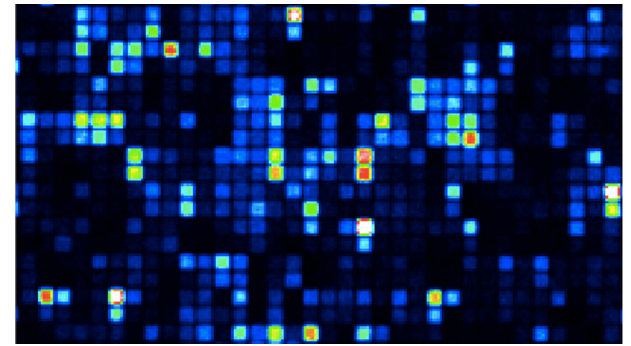
# Genentech



- Develops drugs to treat significant medical conditions
  - mainly cancer, but also Alzheimer, Parkinson, flu, hepatitis, arthritis
- Uses genetic information to produce drugs
  - Need to provide researchers with new genetic strains so as to understand the cause of diseases and to test new drugs.



# Genentech



- Using a relational DB, the Genentech team needed to change the schema every time they introduced a new experiment
- MongoDB is able to capture the variety of data generated by genetic tests and integrate it with the existing Oracle DB
- This reduced development time from months to weeks or even days.

# Conclusion

- Making the right choice of database can deliver quantified business results.
  - Shorter development time
  - Improved customer experience
  - Enabling new types of applications
  - Achieving higher revenues or reducing costs